

# Smart Document Q&A System

## System Architecture Documentation

---


### 1. Overview

The **Smart Document Q&A System** is a production-oriented, multi-agent platform designed to provide accurate, context-aware question answering over large, unstructured document collections. The system combines **hybrid search (keyword + semantic)** with **agent-based orchestration** and **intelligent memory management** to ensure high-quality, explainable responses with citations.

The architecture is modular, scalable, and suitable for enterprise use cases such as healthcare, legal, insurance, and internal knowledge management.

---

### 2. High-Level Architecture



## 3. Core Components

### 3.1 FastAPI Application Layer (app/main.py)

**Responsibilities** - REST API exposure - Request/response validation (Pydantic) - Error handling and status management - CORS configuration - Orchestration of downstream services

**Key Endpoints** - POST /api/v1/documents/upload – Upload and index documents -  
POST /api/v1/ask – Ask questions over indexed documents - POST  
/api/v1/conversations/ – Create/manage conversations - GET /health – Health and  
readiness checks

---

### 3.2 Agent Orchestration Layer (app/agents/)

The system uses **LangGraph** to orchestrate specialized agents with explicit state transitions and traceable reasoning.

#### 3.2.1 Query Analyzer Agent

**File:** query\_analyzer.py

**Purpose** Analyzes the user query to determine intent and optimal retrieval strategy.

**Key Functions** 1. Classifies query intent (factual, conceptual, comparison, procedural)  
2. Selects search strategy (keyword, semantic, hybrid) 3. Reformulates the query if required  
4. Records decisions for observability

#### Sample Output

```
{  
    "agent": "QueryAnalyzer",  
    "strategy": "semantic",  
    "reason": "Conceptual question requiring contextual understanding"  
}
```

---

#### 3.2.2 Search Coordinator Agent

**File:** search\_coordinator.py

**Purpose** Executes hybrid search with dynamically tuned parameters.

**Workflow** 1. Receives analyzed query and intent 2. Adjusts keyword vs semantic weights  
3. Executes search via Typesense 4. Ranks and filters results 5. Prepares source citations

**Dynamic Weighting Logic** - Factual: Keyword 0.7 | Semantic 0.3 - Conceptual:  
Keyword 0.3 | Semantic 0.7 - Balanced: Keyword 0.5 | Semantic 0.5

---

### 3.2.3 Response Generator Agent

**File:** response\_generator.py

**Purpose** Generates grounded, well-structured answers with citations.

**Workflow** 1. Receives ranked search results 2. Optimizes context window 3. Synthesizes information across sources 4. Generates final response with references 5. Handles low-confidence or empty-result scenarios

---

## 3.3 Hybrid Search Engine (app/search/hybrid\_search.py)

**Technology Stack** - Typesense - OpenAI Embeddings (text-embedding-3-small, 1536 dimensions)

**Search Modes - Keyword (BM25):** Precise term matching - **Semantic (Vector):** Conceptual similarity - **Hybrid:** Weighted combination of both

### Hybrid Scoring Formula

Score =  $\alpha \times \text{Semantic\_Score} + (1 - \alpha) \times \text{Keyword\_Score}$

### Index Schema

```
{  
    "id": "chunk_uuid",  
    "document_id": "doc_uuid",  
    "document_name": "filename",  
    "content": "text",  
    "chunk_index": 0,  
    "page_number": null,  
    "embedding": [1536 floats]  
}
```

---

## 3.4 Memory & Context Management (app/memory/)

### 3.4.1 Conversation Memory

**Purpose** Persistent storage of conversations, messages, and search history.

**Key Capabilities** - Conversation lifecycle management - Message-level metadata and token tracking - Search strategy logging - Analytics-ready schema

**Database Tables** - conversations - conversation\_messages - search\_history

---

### *3.4.2 Context Manager*

**Purpose** Maintains an optimal context window under token constraints.

**Strategy** 1. Preserve most recent messages 2. Include older messages if space allows  
3. Summarize older context using LLMs when required 4. Enforce token budget limits

**Benefits** - Prevents context loss - Avoids token overflow - Improves response consistency

---

## 3.5 Document Processing ([app/document\\_processor.py](#))

**Purpose** Transforms raw documents into searchable, semantically meaningful chunks.

**Chunking Configuration** - Chunk size: 500 characters - Overlap: 50 characters - Sentence-aware boundaries

**Supported Formats** - TXT - Markdown - Extensible to PDF and DOCX

---

## 4. End-to-End Data Flow

### 4.1 Document Ingestion

1. User uploads document
2. API validates format
3. Document is chunked
4. Metadata stored in DB
5. Chunks embedded and indexed
6. Success response returned

### 4.2 Question Answering

1. User submits question
  2. Conversation context retrieved
  3. Agent workflow executed
  4. Hybrid search performed
  5. Answer generated with citations
  6. Conversation updated and stored
- 

## 5. Key Design Decisions

### Agent-Based Architecture

- Clear separation of responsibilities
- Explainable and debuggable decisions

- Easier future extensibility

## Technology Choices

- **LangGraph:** Deterministic agent workflows
  - **Typesense:** Native hybrid search with low operational overhead
  - **SQLite:** Lightweight persistence, easy migration path
- 

## 6. Performance Characteristics

**Typical Latency** - Embedding: 100–200 ms - Search: 50–100 ms - LLM Response: 1–3 seconds

**Scalability Options** - Horizontal API scaling - PostgreSQL migration - Typesense clustering - Redis caching layer

---

## 7. Security & Compliance

**Current Measures** - Input validation - ORM-based SQL protection - Environment-based secrets

**Planned Enhancements** - Authentication & authorization - Rate limiting - Encryption at rest and in transit - Audit logging

---

## 8. Future Roadmap

**Short Term** - PDF/DOCX support - Embedding cache - Authentication

**Mid Term** - Feedback-driven re-ranking - Multimodal documents - Advanced summarization

**Long Term** - Autonomous multi-agent collaboration - Distributed deployment - Fine-tuned domain embeddings

---

## 9. Conclusion

This system demonstrates a **production-ready, explainable, and scalable** approach to document-based question answering by combining:

- Agent-driven reasoning
- Hybrid search retrieval
- Persistent memory management
- Enterprise-grade extensibility

It is well-suited for real-world deployments requiring accuracy, traceability, and performance.